

Shoppers do come back. They come back through **search** — and the wishlist behaves as if the save never happened.

19,143 documents said the problem was **fit and trust**. 16 survey responses said the failure happens one step earlier — at **re-entry**. Both are true; the discovery engine was structurally unable to see the second one. The MVP is built for the moment the engine was blind to, and it spends information rather than money.

THE METRIC

$W30 = A \cdot \text{quality of save} \times B \cdot \text{re-entry} \times C \cdot \text{doubt resolution} \times D \cdot \text{mechanics}$. We attack **B-self** $\times$ **C** — the only pair that is unowned *and* permitted by the no-incentive constraint.

THE BUILD

A "From your Wishlist" module that fires only on expressed search intent. Live: 546 tests, 10 measured acceptance gates, 5 experiment arms, a validated 30-day cohort model.

THE THREAD

W30 decomposed → 19,143 docs, 12 themes → survey n=16 + 6 walkthroughs → "the wishlist stores the item and discards the question" → search-reconnection MVP → +15% relative W30, guarded by return rate

BUSINESS METRIC DECOMPOSITION

Four terms, equal leverage — so we pick on **headroom, ownership and the constraint.**

A quality of save × B re-entry within 30d × C doubt resolution × D conversion mechanics = W30
User-level, rolling 30 days from *each item's* save event, reported **gross and net of returns**. A colour swap counts as a conversion (SKU-family match).

A · QUALITY OF SAVE
72.2%
carry genuine purchase intent.
Not our lever — it is what makes the problem worth solving.

B · RE-ENTRY
split in two
B-scheduled (price alerts, push) is owned — and banned for us.
B-self (they search again) — **nothing serves it.**

C · DOUBT RESOLUTION
66.1% → 21%
name a doubt the product never answers; **21%** tried to resolve it and still could not.

D · MECHANICS
24.8%
of checkout-stage need is **price** — the one friction the brief bans us from touching.

THE ARITHMETIC — AND THE ARGUMENT THAT SURVIVES SCRUTINY
 $W30 = 0.722 \times 0.45 \times 0.34 \times 0.85 = 9.4\% \rightarrow C: 0.34 \rightarrow 0.44 \rightarrow 12.2\% (+2.8pp, +29\% \text{ rel.})$
"C has the most leverage" is **arithmetically false** — in a multiplicative model a +25% relative lift on *any* term moves W30 by +25%. The case is **headroom** (C sits near 34%, D near 85%), **ownership** (B-scheduled and D are already optimised), and the fact that the monetary ban removes D's dominant friction outright.

BASELINE · TARGET · POWER	
Modelled baseline	9.4%
Survey-measured direction (Q4, n=14)	14%
Commitment, 8 weeks	+15% rel.
Headroom the model implies	+29% rel.
Per arm at +15% MDE (95%/80%)	7,169

Limit, stated once: the corpus fixes the *mix* inside C, not the absolute level of any term — B and D are assumed pending one day of internal instrumentation. **+15% is the commitment because that is what the sample can detect; +29% is stated as headroom, never as the promise.**

Six sources, five stages, four models — and an audit that changed the answer.

0 · INGEST 19,143 unique docs after dedupe, of 19,174 raw. Apr 2012 → Aug 2026.	1 · RELEVANCE 3,922 classified relevant (20.5%). 0 unjudged — every document got a verdict.	2 · TAXONOMY 12 themes induced from the corpus, then hand-edited three times. Not a preset list.	3 · TAG 3,922 theme, severity, journey stage, intent, segment. 87.5% carry a verified verbatim quote.	4 · SCORE $\sqrt{\text{reach}} \times \text{sev}$ $\times \text{metric-proximity} \times \text{tractability}$, with Wilson intervals on every reach figure.
---	--	--	--	---

SOURCE MIX — AND WHY IT LOOKS LIKE THIS

Source	Classified	Relevant	Rate
YouTube comments	7,705	2,144	27.8%
Sitejabber	153	37	24.2%
Reddit	1,632	338	20.7%
Competitor apps	4,571	905	19.8%
Google Play	3,891	409	10.5%
Apple App Store	1,191	89	7.5%

YouTube supplies **55%** of the relevant corpus and is **3.7×** richer than app-store reviews. Adding it was the highest-leverage decision in the build.

WHAT MAKES THE NUMBERS TRUSTWORTHY

- A **50-document audit** found two real defects. One was fixed — theme 9 re-tagged, 51 documents correctly *lost* their only theme. The second is documented and open, and the audit was model-run: a human spot-check is still outstanding.
- **Wilson confidence intervals** on every reach share, so rank order is separable from noise.
- **Tag stability is 63%** on themes — shares are directional at the point and firm in their ranking. Said once, here.

THE QUERY LESSON, KEPT BECAUSE IT COST MONEY

Reddit queries containing "myntra" returned **26.1%** relevant. Ten "gap" queries about occasion, styling and bookmarking that did not returned **9.2%**, and pulled in r/AmltheAsshole. **\$1.20 for 48 usable documents**. A corpus only tells you about the words you went looking for.

The engine ranked **trust and fit** first — and put the two themes we ended up building on **9th and 12th**.

#	Theme (Myntra-only, n = 3,017)	Reach	Sev	Score
1	Quality, Authenticity & Seller Trust	23.3%	4.26	0.234
2	Unreliable Size & Fit Info	19.5%	3.25	0.169
3	Cumbersome Return & Exchange	9.8%	4.55	0.150
4	Misleading Visual Media	7.9%	3.49	0.105
5	Missing Product Specifications	8.6%	2.90	0.071
6–8	stock signals · pre-purchase support · price volatility			
9	Wishlist Interface Friction	0.89%	3.19	0.037
10–11	absent social proof · delivery timelines			
12	Passive Wishlist, No Re-engagement	0.33%	2.80	0.011

HOW THE RANK IS BUILT — AND WHAT IT IS BLIND TO

Score = $\sqrt{\text{reach} \times \text{severity} \times \text{metric-proximity} \times \text{tractability}}$. Trust leads on reach (704 documents) and severity (4.26 of 5). **Ranks 9 and 12 lose on reach alone** — Wishlist Interface Friction in fact carries the **highest tractability in the whole set (0.90)**. And reach here is the share of people who *wrote something*, which is the assumption slide 5 breaks.

INSIDE THE 66.1% — WHAT SHOPPERS SAY THEY COULD NOT DETERMINE

Will it fit me?		46.2%
Is the fabric what it claims?		17.3%
Can I trust this seller?		14.7%
Will it look like the photos?		13.3%
Price · returns · social · styling		11.3%

Deliberation concentrates at **evaluation (67.9%)**, not checkout (3.9%). Nothing is within 28 points of fit.

WHAT WE DROPPED, AND WHY THAT MATTERS

Return anxiety read as a top-three blocker at **11.3%** full-corpus. Restricted to Myntra and to pre-purchase it is **1.9%** — returns are what people complain about *after*, not what stops them buying. We dropped it. Fit moved the other way on the same correction: **21.1% → 30.6%** of pre-purchase documents, and 46.2% of named doubts.

The engine measured complaint volume; primary research asked about the **moment of return**.

CONFIRMED

- **Fit-adjacent doubt is real and broad.** Fabric, photos, "suit my body", fit — together they touch **8 of 14** (57%) wishlists (Q12).
- **Doubt genuinely goes unresolved.** **21%** answered *"I tried, and still couldn't find out"* (Q8) — the first direct measurement of term C, which the corpus could not produce.
- **Two doubts, not one population.** Fit×trust overlap is **1 of 14** — "one shopper, mislabelled" is not supported.

CHALLENGED

- **Asked for one blocker, shoppers do not say fit.** Price/sale takes **7 of 14** (50%); fit + trust + fabric take **4 of 14** (29%) (Q7, single-select).
- **A blocker the corpus never surfaced:** *"I'd forgotten about it until now."*
- **The failure is shaped like comparison, not ignorance.** **36%** buy none of the close alternatives or leave to find more; **14%** buy none outright (Q11).
- **W30 finally has a measured direction:** **14%** bought from the wishlist in 30 days (Q4).

WHY BOTH RESULTS ARE TRUE — WHAT DIFFERS IS THE INSTRUMENT, NOT THE FINDING

The corpus is **reviews, comments and posts** — written *after delivery*, about the object. It sees a doubt **once, never twice**, and **1.56%** of it mentions a wishlist. "Passive Wishlist" scored 12th on 0.33% reach not because the behaviour is rare, but because **nobody** narrates it.

The survey asked about a **specific saved item, now**. At that moment price salience wins the single-select, fit-adjacent doubt is broader than any one label, and the dominant failure is a decision left open. **Term C is the content of the doubt; B-self is whether it is ever re-opened.** The engine ranked C well and could not see B at all.

So the MVP does not choose between them. It fires at self-initiated re-entry (B-self) and carries the saved decision back with enough evidence to close it (C) — the one place both terms are addressable with no discount and no notification.

Method and limits, stated once: [survey instrument and coding rules published with the dashboard](#) · 16 responses → 15 valid shoppers → **14 item-level**, plus 6 moderated walkthroughs. At n=14 one respondent moves any share ~7pp — directional confirmation of corpus-scale findings, not proof.

PROBLEM DEFINITION

The wishlist stores the item and discards the question — so the shopper carries the risk.

The ladder — three measured rungs, two under test		Evidence
1	Wishlisted items do not get bought	brief
2	"I'm not sure it'll fit" · "I'll come back to it"	measured
3	The wishlist stores the SKU and throws the question away	teardown
4	What would settle it is never joined to <i>this</i> shopper when they ask again	Q8 · 21%
5	The shopper carries the entire risk of being wrong	walkthroughs

WORKAROUNDS — LEVEL 5 OBSERVED, NOT ARGUED

What they do	What it costs them
Screenshot into WhatsApp for an opinion	a favour, and the wait
Hunt reviews for a similar body	time, often fruitless
Order two sizes, return one	money up front
A notes app as the real shortlist	duplicated effort

Read the right-hand column: the platform has **externalised the cost of resolving doubt onto the shopper, and they are paying it**. Ordering two sizes is buying insurance from Myntra, which books it as a return.

SEGMENT — WHERE THE SIGNAL CONCENTRATES
Fit-uncertain shoppers — 28.3% of documents, a **3.99×** theme lift (the strongest in the corpus), 84.5% blocked on fit and size. **New / low-trust** — 20.9%, 2.81× lift, 55.7% on seller trust. Shares of *documents*, not users.


THE SURFACE TODAY — TEARDOWN
A flat reverse-chronological grid with no memory of *why* anything was saved — plus a price-drop alert, a discount badge, a cashback banner and a coupon. **Every lever is monetary.**

WHY THE HARDEST CONSTRAINT IS ACTUALLY THE INSTRUCTION
A discount does not reduce risk — **it prices it and leaves it exactly where it was**. Every lever on this surface today is monetary, and every one leaves the item sitting there. Cash on Delivery is the precedent, for the framing not the feature: framed as "low card penetration" it buys a payments integration; framed as "the customer carries all the risk upfront" it buys COD. **Re-allocate the risk rather than discount the price of bearing it.**

Falsification, pre-committed: if the walkthroughs describe *not knowing* rather than *not wanting to be wrong*, level 5 is unsupported — we keep level 3, drop the risk framing, and say which read produced that. A workaround named by 4 of 6 goes on a slide; named once, it stays in the appendix.


Search is where the decision re-opens, so the wishlist meets them there.

1 · "DID I SAVE THIS?"




Fires only on expressed search intent. Saved colour and size, two co-equal actions.

2 · "STILL RIGHT FOR ME?"



Ten signals, each citing a source. Fit reads INFO, never a score.

3 · "WHAT ELSE MATTERS?"



A priority reorders and never filters. No winner is named.

CONSTRAINTS, ENFORCED AT LINT LEVEL

C-1 no discount, coupon, cashback or sale urgency · C-2 no push in v1 · C-3 search renders independently, fail-open in ≤ 250 ms · C-4 precision > recall · C-6 logged-out sessions reveal nothing · C-7 accessibility is a launch gate.

WHAT WE DELIBERATELY DID NOT BUILD

A fit score — permanently status: "unknown". No size chart or fit feedback exists here, so a score would be invented data wearing the costume of evidence. Also unbuilt and named on the surface: review *quality* as distinct from count, seller trust, stock depth, visual similarity.

546
automated tests

10 / 10
acceptance gates pass

0
silent substitutions,
10,000-run fuzz

0
leaks, 66 logged-out
calls

+0.02 ms
search cost, vs 120 ms
budget

Live prototype: wishlist-reconnection-prototype.vercel.app — screens captured from the current build; all 10 module states reachable through the built-in harness, verified logged-out. 5 experiment arms, kill switch armed.

SUCCESS METRICS

The metric that proves the judgement is a **guardrail**, not the north star.

Metric	Definition, and why this one	Term	Baseline	Target
North star · W30	buyers of a previously-wishlisted item within 30d of its save ÷ wishlisters, gross <i>and</i> net of returns. Net is honest; gross says why net moved.	all	9.4% modelled; 14% survey (Q4)	+15% rel. in 8 weeks
L1 · Opportunity rate L2 · Module engagement	searches with ≥1 matching saved item ÷ searches taps ÷ renders. Without L2 a flat L3 is unreadable — a bad answer looks like an unseen one.	B-self	41.9% shadow 0 — no surface	monitor ≥ 40%
L3 · Doubt-resolution rate ★	items reaching a terminal state — add-to-cart or explicit remove — within 7d of view ÷ items viewed. Add-to-cart alone scores whether it <i>sells</i> — what a discount does. This scores whether it <i>decides</i> .	C	34% modelled; 66.1% carry a doubt	34% → 44%
L4 · Decision latency L5 · Comparison completion	median days save → terminal comparisons ending in a decision ÷ opened — the failure mode research found, and Compare serves.	C	unknown Q11: 36% fail	falls monitor
G1 · Return rate on assisted orders ★	returns ÷ orders where the module was viewed before add-to-cart. Confident wrong-size buys raise gross W30 while destroying net value, and read as a win.	—	control	≤ control, no tolerance
G2 save rate · G3 dismissal rate G4 AOV & units · G5 render delta	G3 catches the quiet failure — a module that trains shoppers to ignore it. G4 guards the cheap win of steering people to low-value items.	A, B, D	control	≥−2% rel. · falling flat · ≈ 0

The counter-metric, named before anyone asks: remove/archive rate **will rise, and that is a success** — closing a decision negatively still closes it, it raises the signal density of the list, and it is why the positive answers get believed.

Experiment & decision rules: 5 arms randomised on wishlisters (W30 is user-level); ~7,200/arm for a +15% MDE at 95%/80%; 7–8 weeks. Read with a **confidence sequence, not a fixed horizon** — peeking 40× inflates fixed-horizon false positives to **33.7%** against a nominal 5%; the sequential bound holds at **2.3%**. gross↑ net↑ G1 flat → **ship**; gross↑ net flat G1↑ → **fail, do not ship on the gross number**.

RISKS & MITIGATION

The largest risk is not that the feature fails. It is that the experiment cannot tell us whether it did.

Risk	Why it is real — measured, not feared	Mitigation, pre-committed
R1 · The panel cannot answer the question	A 400-tester panel supports only a 70.66pp lift. The real 5pp lift needs 19,390; the B–A question needs 1,22,787. Five arms, intent-to-treat dilution, a between-treatment contrast.	Re-scope the panel to what it can answer: comprehension, recovery from an unavailable variant, whether the 10 states read. Run the powered A/B on real traffic only, and power for B–A or drop the B arm.
R2 · The evidence is thinnest where the decision was made	The corpus cannot see re-entry — 1.56% of it mentions a wishlist. The instrument that can is n=14 item-level, where one respondent moves a share ~7pp.	Falsification tallies fixed before the sessions: a workaround named by 4 of 6 earns a slide; if mostly_waiting wins, the problem is timing and the direction changes.
R3 · A false positive trains shoppers to ignore the module	A wrong "you saved this" is worse than a miss, and unrecoverable — the next correct one is not believed either.	Exact matching only, no embeddings in v1. Thresholds keyed per modality (0.72 text / 0.82 voice / 0.85 image). Below threshold: render nothing. Dismissal suppresses and logs a signal.
R4 · Our own gate numbers flatter us	Precision labels are generator-derived, and the generator infers brand and type the way the matcher does — a shared error is invisible. Test queries carry no typos.	Quoted as regression guards, not launch evidence. Before ramp: a hand-labelled 500-pair set, queries sampled from real traffic, a 500 rps load test.
R5 · Honest answers cost short-term GMV	Remove/archive rises by design, and resolving a doubt toward "no" books zero revenue this quarter.	Measured as 60/90-day repeat purchase rate in the treatment cell, with the direction declared in advance so it cannot be renegotiated at read-out.
R6 · We contradicted our own pre-registered rule	Our guardrail said "this corpus does not support a wishlist-UI-centric MVP." The shipped MVP is one, and fit in it is permanently unknown.	Amended with the reasoning, dated — not quietly dropped. The rule was right about the corpus and wrong about its coverage.

Privacy is a launch blocker, not a risk row: wishlist data never crosses the account boundary. Gate: 0 wishlist fields across 66 unauthenticated calls, covering response payloads and the shadow topic — timing side-channels are not covered and need their own pass before ramp.

IN CLOSING

The engine found the doubt. Shoppers found the moment. The build is instrumented to be proved wrong.

BUSINESS METRIC

W30 decomposed into four terms. **B-self** and **C** are the only unowned pair the no-incentive constraint leaves standing.

AI DISCOVERY

19,143 docs → 3,922 coded → 12 themes. **66.1%** name an unanswered doubt; **46.2%** of those are fit.

PRIMARY RESEARCH

n=16 plus 6 walkthroughs. Confirmed the doubt, **overturned the moment**, and exposed the corpus's blind spot.

PROBLEM

The wishlist stores the item and discards the question — and the shopper carries the risk of being wrong.

MVP

Search-reconnection, live. Memory plus agency; no money, no notification.

WHAT IS ACTUALLY BUILT AND RUNNING

- **Discovery engine**, deployed with a live cohort toggle and a working classify endpoint — [dashboard](#)
- **Prototype**, 546 tests, 10 measured gates, all 10 module states reachable — [live link](#)
- **The measurement pipeline is itself validated**: the staged ramp held assignment across 13,182 users with 0 arm changes, the kill switch fired on a 74.5% guardrail breach, and sequential inference held its error rate where fixed-horizon peeking did not.

WHAT WOULD CHANGE MY MIND

If the walkthroughs describe **not knowing** rather than **not wanting to be wrong**, the risk framing goes and level 3 stands alone. If `mostly_waiting` wins the tally, the problem is timing and this is the wrong build. If L2 is high and L3 is flat, the module is seen and useless — that is a kill, not an iteration.

In one line: a save is an open question, and the only non-monetary moment to close it is the next time the shopper asks it themselves.

Limits carried to the end, not buried: no absolute term baselines until day 1 of internal instrumentation · n=14 item-level survey · gate numbers are generated-label regression guards · no real fit score, review-quality or seller-trust signal, and no visual similarity · the recruited panel answers usability, not lift.