

01 THE OPPORTUNITY

Shoppers Do Come Back.

They Come Back Through Search.

Grow saved-item conversion within 30 days by protecting gross margins, avoiding promotional dilution, and driving organic conversion.

THE MARKET CHECK

51.5M

weekly actives — India's largest fashion platform

CLSA, 2026

₹6,043cr

FY25 revenue, up 18% year on year

MCA filings

18x

profit growth on the same demand

MCA filings

9.4%

of wishlisters buy a saved item within 30 days

= 0.722 × 0.45 × 0.34 × 0.85 (A×B×C×D). Intent measured; B, C, D assumed until day-one instrumentation.

Myntra grows by converting shoppers it already has, not by buying new ones.

Demand is already paid for.

THE PREMISE

Four things have to go right for a saved item to sell. Fix any one and the whole number moves — so we picked the ones nobody owns.

A

They meant to buy it — 7 in 10 saves are genuine

near ceiling

B

Do they come back? Alerts exist already and cost margin. When they search instead, nothing helps.

unowned

C

Do they get an answer? 2 in 3 have a doubt we never answer. 1 in 5 who looked gave up.

unowned

D

Do they buy? Price stops about 1 in 4 — and price is off-limits

excluded lever

A and D are already handled. B and C are wide open.

THE STRATEGIC GOAL

Take W30 from 9.4% to 10.8% in eight weeks — +15% relative — on organic conversion, with no discount, coupon or cashback spend behind it.

Why it matters:

For shoppers — the saved item comes back unprompted, and the doubt that stopped them is answered in the same place.

For Myntra — on a 10M-wishlister month, +1.4pp is ~140,000 more closed decisions, about ₹28 Cr of GMV at a ₹2,000 basket.

For margin — every point of that lift is earned without promotional dilution, on demand already paid for.

9.4% → 10.8% · ≈₹28 Cr/mo · zero discount spend

02 HOW THE DISCOVERY ENGINE WORKS

Live: Discovery Engine ↗

Turning Customer Voice Into Actionable Insight, One Comment at a Time.


Six public sources, five steps — and an audit that changed our own answer.

- 1

Reading **19,143** comments by hand is slow and inconsistent.
- 2

One model reads every comment the same way, so nothing is missed or judged differently.
- 3

Loose complaints become ranked, quotable evidence.



03 WHAT THE ENGINE FOUND

The Data Speaks: Trust and Fit Win.

Our Feature Ranks Ninth and Twelfth.

What 19,143 comments settled, what they could not — and the four claims left for primary research to test.

RANKED BLOCKERS · SHARE WHO RAISE IT

Across 3,017 Myntra comments about a buying decision.

Quality & seller trust		23%
Unreliable size & fit		20%
Difficult returns		10%
Missing product details		9%
Misleading photos		8%
The wishlist itself two themes, 9th and 12th		1%

Returns hurt badly when they happen — but hurting is not the same as stopping the sale.

ONE BARRIER, NOT FOUR

The doubts people actually name:

Will it fit me?		46%
Is the fabric right?		17%
Can I trust the seller?		15%
Will the colour match?		13%

One failure, not four: 9 in 10 doubts are things the page cannot verify.

Two in three hit a barrier they cannot resolve.

WHAT THE ENGINE CANNOT SETTLE

Does it stop the sale, or only delay it?

What brings someone back to a saved item?

Would answering it change the decision?

Fewer than 2 in 100 comments mention a wishlist — reviews are written after delivery. Primary research had to settle these, as four claims with tests.

CLAIM 1

Fit uncertainty is the single largest blocker.

Size & fit · 46% of named doubts

THE TEST

Disproved if shoppers rank price above fit when asked.

CLAIM 2

The doubt cannot be resolved on the product page.

Blocked before buying: 2 in 3

THE TEST

Disproved if the answer was on the page and they missed it.

CLAIM 3

Returns are a post-purchase grievance, not a blocker.

Ranked 3rd on raw counts only

THE TEST

Disproved if returns still rank once the split is applied.

CLAIM 4

The wishlist itself is not what people complain about.

Two wishlist themes, both near the bottom of twelve

THE TEST

Disproved if wishlist usability is named in research.

Where Data Meets Human Truth: Naming the Moment of Doubt.

Primary research: 16 survey responses and 11 prototype testers across two rounds — a method that never asks anyone to hand over a real wishlist.

✓ WHAT IT CONFIRMED

- ✓ Fit-type worries are real and widespread — fabric, photos, "will it suit me", sizing: **more than half** the lists we looked at.
- ✓ The doubt really does go unanswered. Of those who went looking for an answer, **1 in 5 came back empty-handed**.

✗ WHAT IT CHANGED

- ✗ Asked for their single biggest blocker, **half said price or a sale** — not fit.
- ✗ A reason the data never surfaced: *"I'd forgotten about it until now."*
- ✗ **More than a third** get stuck comparing. And **four of eleven** prototype testers named a blocker our data has no word for — *can't tell them apart, might be out of style, don't know what it goes with*. Not a missing fact. An unfinished decision.

01 · SAVE

The shopper saves something because it might matter later.

02 · SEARCH AGAIN

Days later they search again, carrying none of that context.

03 · RECONNECT

The step nothing does today — the saved item should reappear.

04 · RESUME

The decision resumes instead of restarting.

WHY BOTH ANSWERS ARE RIGHT

Fewer than 2 in 100 comments mention a wishlist — reviews are written after delivery, so nobody writes about it. Silence is not evidence it works.

Ask about a saved item right now and the blocker isn't missing information — it is a decision left open and never picked back up.

05 THE PROBLEM, STATED

The Wishlist Remembers What You Saved. It Forgets Why You Cared.

The wishlist answers none of the four doubts that stop a saved purchase. Segment: fit-uncertain & low-trust shoppers (48%).

⚖ THE PROBLEM, IN NUMBERS

W30 near 9% — nine in ten saved items go unbought inside 30 days.
2 in 3 name a doubt we never answer; 1 in 5 who looked came back empty.
Every lever there is monetary — each leaves the item sitting.

FIVE WHYS — METRIC TO ROOT CAUSE

- 1

Saved items don't get bought

objective
- 2

"not sure it'll fit" · "can't tell them apart" · "I'll come back to it"

measured
- 3

The wishlist saves the item, not the question

measured
- 4

What would settle it never reaches this shopper

survey
- 5

So the shopper carries the risk of being wrong

to test

THE VALUE TO THE SHOPPER

They screenshot into WhatsApp, hunt reviews, keep the real shortlist in Notes — and order two sizes to return one, money up front. We handed them the cost of their own doubt.

WHAT THE RESEARCH PRODUCED

- 19,143 comments

discovery engine

Twelve blockers ranked: trust 23%, fit 20%. The wishlist itself ranks 9th and 12th — fewer than 2 in 100 comments mention one at all.
- 16 survey responses

primary research

Half name price or a sale, more than a third get stuck comparing, and one reason the corpus never surfaced: "I'd forgotten about it until now."
- 11 prototype testers

two rounds, live build

All eleven found the saved item; five still could not say why it appeared. Four named a blocker the corpus has no word for.

HOW THE THINKING MOVED, STAGE BY STAGE

- Business metric

→ four terms; two unowned
- Product outcomes

→ they re-enter alone, the doubt gets answered
- AI discovery

→ 19,143 comments, blind to re-entry
- Primary research

→ doubt confirmed, moment moved
- Problem definition

→ item kept, question discarded

★ THE TAKEAWAY

Not buying isn't indecision. When the shopper carries all the risk, waiting is the sensible move.
If the prototype study shows people simply didn't know rather than didn't want to be wrong, step 5 is wrong — we keep step 3, drop the risk framing, and say so.



06 PRIORITISATION

Eighteen Ideas Went In. One Came Out.

And It Wasn't a Wishlist Feature.

Four questions decided it: does our own evidence back it, could only Myntra build it, can we build it in the time, can we *show* it working. Reach was the same for all three, so RICE would have settled nothing.

18

ideas written down,
all scored the same way

→

3

worth building
— the three below

→

1

built, tested
and live today

The other fifteen never reached this deck. They were too small to matter, too slow to build, or they leaned on a discount — the one lever we were told not to pull.

DIRECTION 1

PARTLY USED

Make the wishlist worth opening.

Split the saved items into decided and still-unsure, answer fit on the card itself, and put near-identical items side by side.

THE CATCH

Every version of it waits for the shopper to open the wishlist. Almost nobody does — a better list nobody opens changes nothing.

DIRECTION 2

SET ASIDE

Answer "will it fit me?" with other people's returns.

Find shoppers who kept the same sizes as you, and say so on the item: *"people with your history kept this one."*

THE CATCH

The best use of data only Myntra holds — but it answers one doubt and leaves the others: comparing, going out of style, what it goes with.

DIRECTION 3

BUILT

Bring the saved item back when they search.

It reappears as they type, carrying the checks that settle the doubt. No reminder, no discount, no trip to the wishlist.

WHY IT WON

It is the only one that starts on its own, without the shopper remembering — and the only one that answers the doubt in the moment it is felt.

Nothing that lost was thrown away. Two of the ideas we ranked below the winner are now parts of it — they became the confidence panel and the compare screen on the next slide. Only the fit idea was set aside, and it is first in the queue after launch.


07 WHAT WE BUILT

Live: Prototype ↗

Reuniting Shoppers With What They Love Before They Even Ask.


The search-interruption prototype. The two screens that carry it — the rest, and what testers made of them, follow.

1 · INTERRUPTION



The one that matters most. Saved items return before the search runs.

2 · RECONNECTION



Three items at most, two equal paths, ranking unchanged.

WHY THE FIRST SCREEN CARRIES THE PRODUCT

Every screen after it assumes the shopper remembers saving the item. **Typing is the only moment where that's still an open question** — and it costs nothing. Saved items sit in a labelled group *above* the suggestions, colour and size spelled out. **Not certain it's a match? The group doesn't appear.**

THE RULES WE HELD OURSELVES TO

No discounts or urgency. No notifications. Search never waits for us. When in doubt, show nothing. Accessibility blocks launch. **What we chose not to build:** a fit score, a review verdict, a seller-trust rating, stock levels, image search. **The data can't support them, and we say so.**

★ THE TAKEAWAY

The interruption is the product. Everything after it is the shopper finishing a decision they'd already started. 546 automated tests, all ten acceptance checks passing, no wrong-variant substitutions in a 10,000-run stress test, and no measurable cost to search speed. ↗


08 WHAT TESTERS SAID

Live: Prototype Survey ↗

A Promising First Step: Users Found It, Trust Is Still a Work in Progress.


The other three screens, and what eleven testers made of each. Counts, not percentages — a direction, not a result.

3 · CONFIDENCE




Ten checks, each naming its source. **Six of eleven** could say why the item appeared.

4 · COMPARISON



Reordering by what matters picks no winner. **All nine who compared** said it helped.

5 · STYLING



Four at most, all already saved. **Nine of eleven** would wear it; **one never found it**.

✓ WHAT HELD UP

✓ All eleven found a wishlisted item while searching — ten right away.

✗ WHAT DIDN'T

✗ Deciding in place got worse — 3 of 5, then 1 of 6.

✗ A real defect: a removed item came back, still labelled saved.

11/11 found it	6/11 could say why	4/11 decided in place	5/11 want an alert
----------------	--------------------	-----------------------	--------------------

WHAT'S STOPPING THEM — ALL 11 BLOCKERS

Price — the excluded lever

4

Doubts we have no word for

4

Fit, quality, or nothing

3

09 HOW WE'LL KNOW IT WORKED

One Number We Mean to Move.

Four That Must Not Move With It.

8-week user-level A/B — Control: search vs Treatment: reconnection module — ~7,200 users/arm (95/80% power, +15% MDE).

★ NORTH STAR METRIC

W30 — wishlist-to-purchase conversion = wishlisters who buy at least one saved item within 30 days ÷ all wishlisters.
Bought = counted **after returns** — the honest number. · **Wishlist-to-bag** (item-level, ~15% benchmark) is a **diagnostic**, not the north star.

▶ LEADING INDICATORS

They move first. If these stay flat, the north star never will.

Module reach and use — does the saved item appear in the search, and is it tapped? Without this, a flat result is unreadable.

Decisions closed — bought or deliberately removed ÷ all saved items. Counts deciding, not just selling.

Time a decision stays open — days from save to a resolution either way. Three days beats thirty.

⚠ GUARDRAIL METRICS

These must not break, even if the north star rises.

Return rate on helped items — a wrong-size push destroys the value behind the win.

Dismissal rate — catches a module people quietly learn to ignore.

Basket size — catches the cheap win of steering shoppers to cheaper things.

Search latency — core search must not slow; the match has a **250ms** hard timeout and fails open.

🎯 TARGETS — FIRST EIGHT WEEKS

North star — W30, +15% relative

today		9.4%
8 weeks		10.8%

Leading indicators
Tapped by **2 in 5** · decisions closed **1 in 3** → **4 in 10** · time falling.

Guardrails
Returns **no higher than normal** · dismissals, basket, speed unchanged.
9.4% is modelled, not measured; external wishlist-to-purchase benchmarks run 5–20%.

★ THE TAKEAWAY

Passing our own tests proves the feature works as built. It proves nothing about whether shoppers buy more.

One result we've committed to in advance: **deliberate removals will go up, and that is success** — closing a decision the other way still closes it. If purchases rise but returns rise with them, we don't ship.



10 RISKS, ROLLOUT AND WHAT'S NEXT

Charting the Path Forward: Embracing Risk, Demanding Clarity.

Four ways this fails, the response agreed in advance for each, and the three phases that would prove it.

 FAILURE MODES & MITIGATION

The test group is too small **biggest risk**

Risk — 300–500 testers; the two-arm read needs ~14,000, isolating *why* needs ~120,000. **Mitigation** — they test comprehension; the read runs live.

One wrong match and the module is dead

Risk — a false "you saved this" is unrecoverable; the next correct one isn't believed. **Mitigation** — exact matching only, nothing shown below the bar.

The module slows core search

Risk — latency added to search costs more than this feature earns. **Mitigation** — matching runs beside search, never before it: +0.03ms measured, 250ms fail-open.

A shared device exposes someone's saved items

Risk — reconnection on a logged-out or shared session leaks a private list. **Mitigation** — signed-in only; signed out, the module shows nothing. Tested, and a launch blocker.

 ROLLOUT — THREE PHASES

Phase 1 · Comprehension **weeks 1–2**

300–500 testers on the live prototype. **Goal** — does it read as "you saved this", does dismissal work, is the removed-item defect gone. Plus a **shadow pass**: matches logged, nothing shown.

Phase 2 · Live A/B **weeks 3–10**

User-level split, ~7,200 per arm, search against the module. **Goal** — W30 9.4% → 10.8%, returns flat.

Phase 3 · Ramp or kill **week 11+**

Ramp only if returns hold. Seen at volume and nothing moves — **a kill, not an iteration.**

WHAT WINNING LOOKS LIKE AT EIGHT WEEKS

Shoppers — the item returns unprompted, doubt answered.

Myntra — ~₹28 Cr/mo GMV, decisions 1 in 3 → 4 in 10.

Instrument first — W30 measured, not modelled. **Fix the removed-item defect** — launch blocker. **Set the kill date** — week 11, in advance.

 THE TAKEAWAY

The wishlist was never the problem. Waiting for someone to open it was.
The limits still stand: 9.4% is modelled, not measured, and passing our own tests proves the build, not the business. Eight weeks and one number settle that — and we wrote down in advance what would make us kill it. 